Visual Analysis of Recurrent Neural Networks

Yao MING

Jan 11, 2017

Introduction

Recurrent Neural Networks (RNNs)

- More general and flexible than Feedforward Neural Network or CNN, operates on **sequential data** like text and audios
- Effective in: language modeling, machine translation, speech recognition, sentiment analysis, image captioning…
- With fast growing research attention

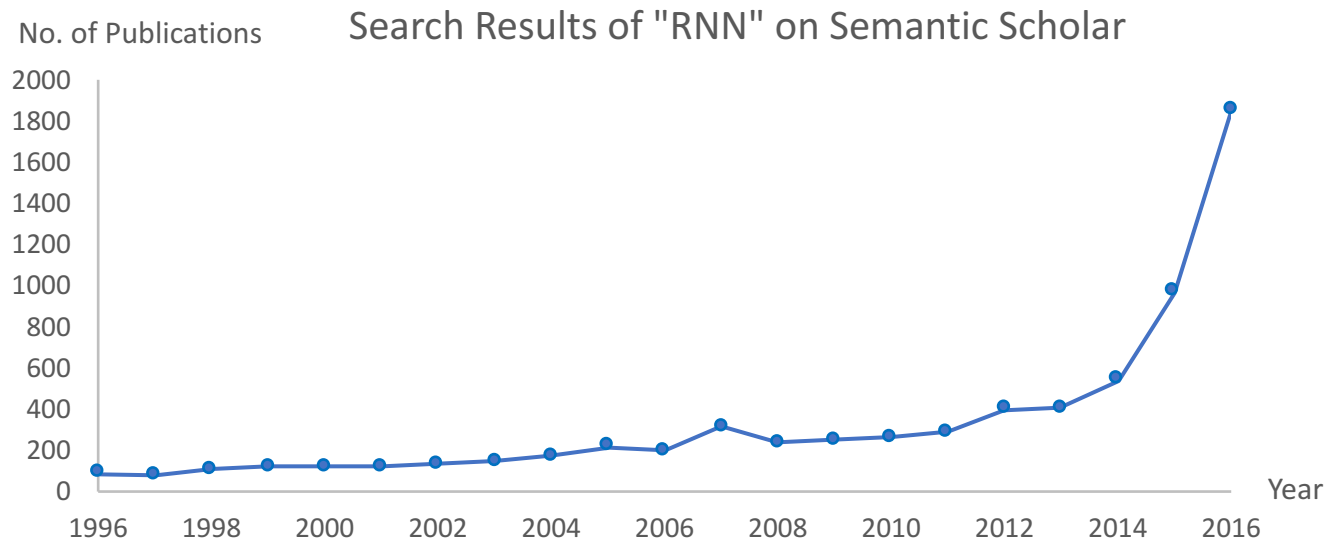

Introduction

Vanilla RNN (single layered)

- Like a scanner, each time step, update **hidden state** s_t using input x_t and previous s_{t-1}, and output o_t.
- Weights are shared over time steps.

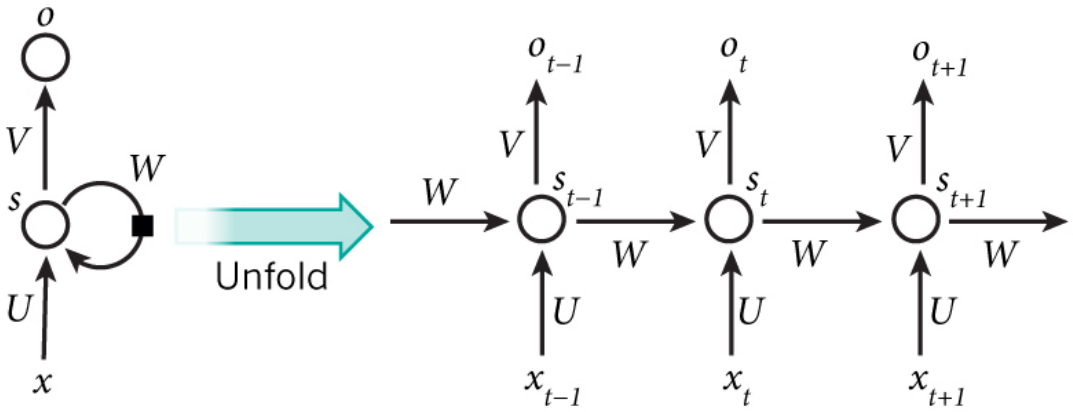

$$s_t = f(Ux_t + Ws_{t-1})$$

$$o_t = softmax(Vs_t)$$

f is nonlinear e.g.: tanh

Variants

- Long Short Term Memory networks (LSTMs)
- Gated Recurrent Unit (GRU)
- Bi-directional LSTMs

Introduction

Word Embedding (compress input of RNNs)

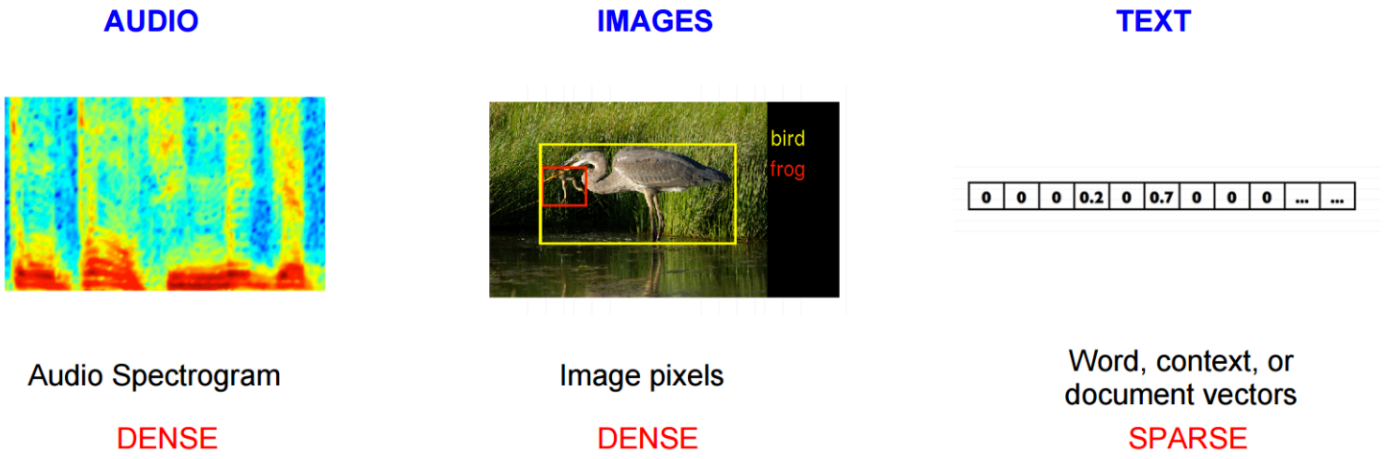

- Compress word vector from **several thousands** of dimensions to **tens or hundreds** of dimensions
- Some of the dimensions may contain semantic meaning or sentiments

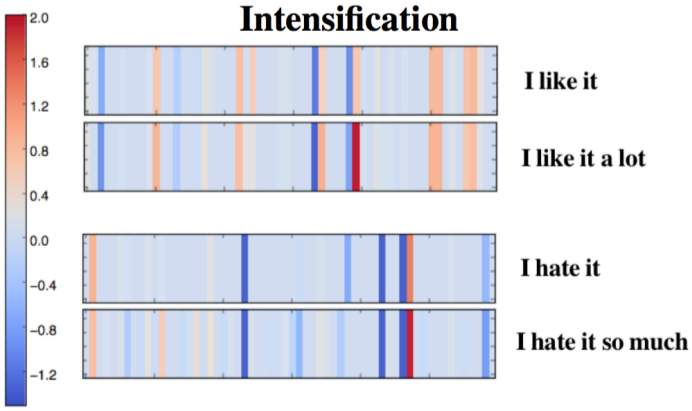

The Problem is

In one sentence,

- How to **compare** different RNNs and better **understand** their **inner mechanisms** (capture keyword, long term memory) to help **improvements** and **understanding** of RNNs

Specifically,

- How to encode the **hidden states** and **gate activation** information into visualization to **intuitively** presents what RNNs learned from text (or audios).
- Design more convenient visualization for **comparatively** analyzing model's **performance/error** on **datasets** for improvements.

The Problem is

An "Old" Problem

- Known issues of interpretability of what deep Neural Networks (with **millions** of parameters) have learned
- Poor knowledge in the **source** of RNNs' impressive performance and the **shortcomings** of different RNNs (Karpathy et. al, 2015)
- Visualizing RNNs are less studied than CNNs, and are more challenging (Images are visualizations, but text/audios are not.)
- Existing work (LSTMVis, 2016) only focus on exploring hidden state patterns of pre-trained models.

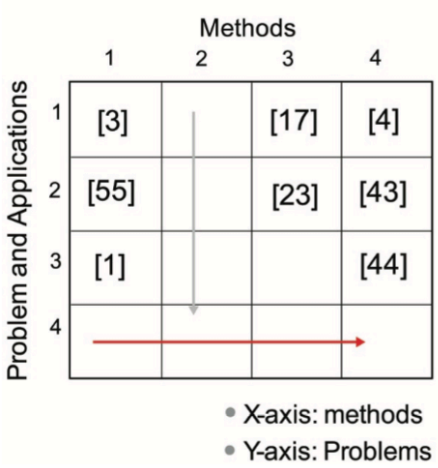

Contributions

- Clearer understanding on RNNs' inner mechanism (long and short term learning ability, hidden state interpretation, gating mechanisms, etc.), which drives further improvements on RNNs' architecture
- A user-friendly visualization tool for:
  - analyzing model's comprehensive performance (debugging)
  - comparing pros and cons of different models (understanding)
  - exploring datasets and alternative models (improvements)

Related Works

Closely Related Visual Analysis of Hidden State Dynamics in Recurrent Neural Networks. (arxiv Jun, 2016)

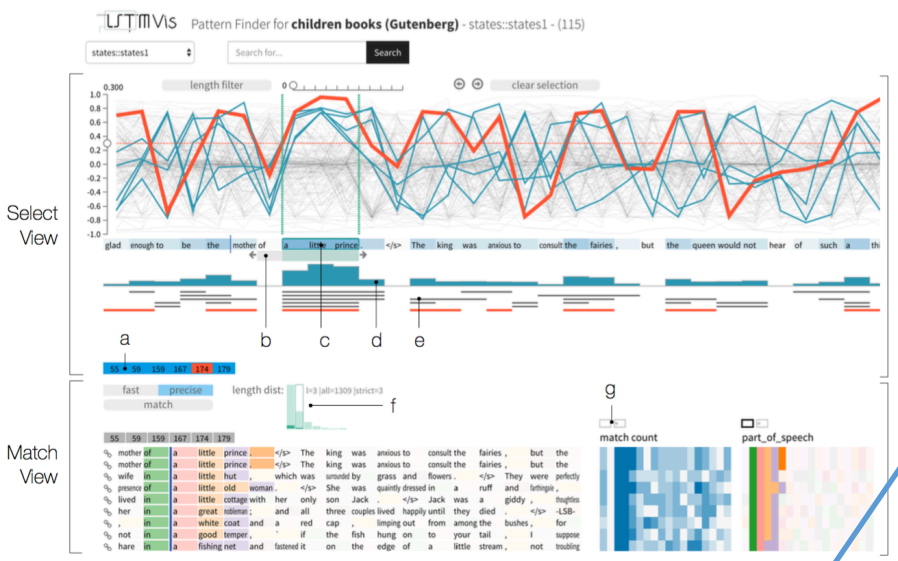

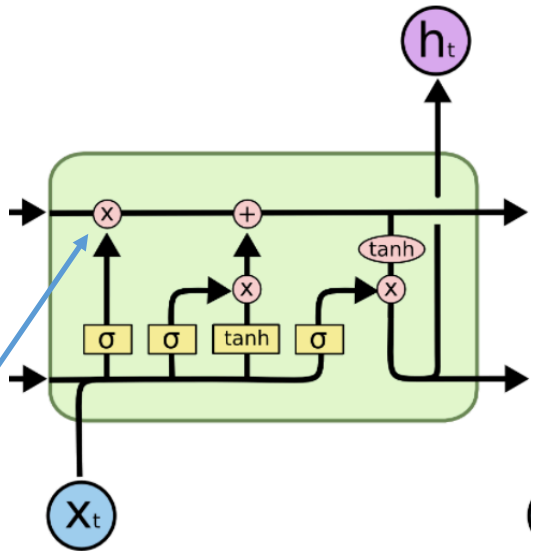

- It fails to analyze **gating structure** -- a **key** design of popular RNNs (LSTM, GRU) to control the update of hidden states and output.
- Also fails to analyze **error/performance** of a RNN (which is important in model evaluation).

Related Works

Understanding RNNs

Visualizing and Understanding Recurrent Networks (Karpathy, A., Johnson, J. and Fei-Fei, L., arxiv 2015 / ICLR 2016):

- a practical **error analysis method** which break errors into different categories to guide model improvements;
- model interpretations which can be of reference in case study.

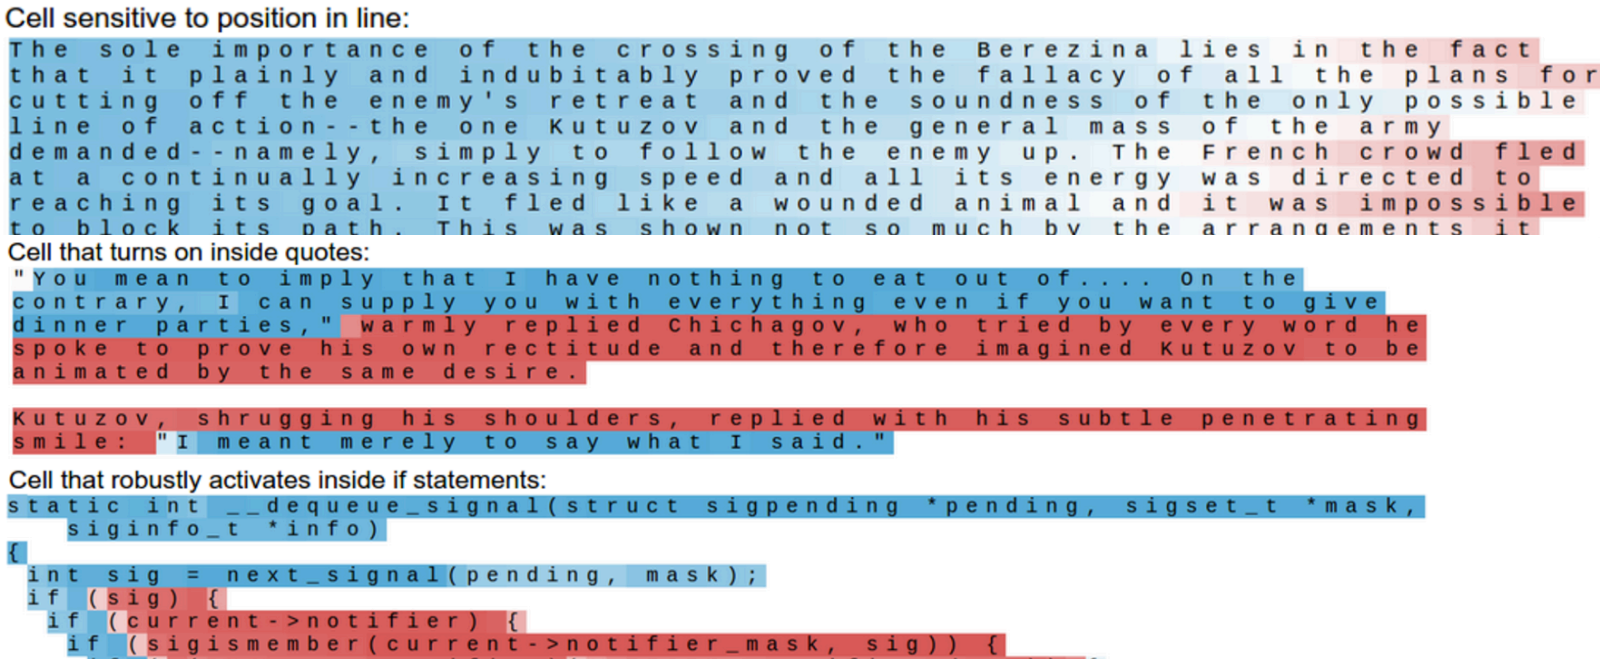

Related Works

Understanding RNNs

Visualizing and Understanding Neural Models in NLP (Li, J. et. al., NAACL-HLT 2016):

- explorations of models' capability in several language phenomenon (concessive sentence, negation, intensification) (for case study)
- visual designs to compare models ability in learning word embedding and capturing keywords (design alternatives)

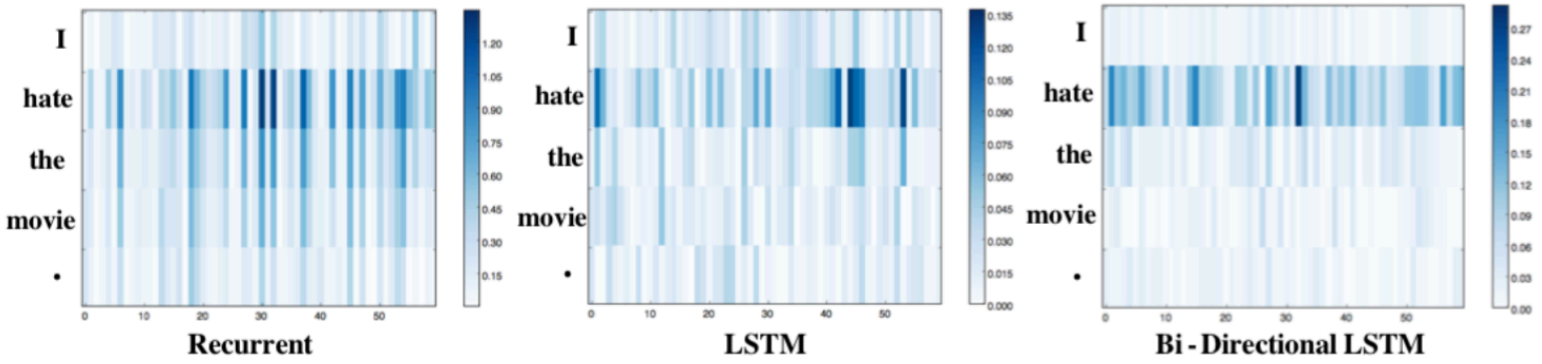

Related Works

Visual analysis of Neural Networks

- Visualizing the Hidden Activity of Artificial Neural Networks (VAST 2016)

  t-SNE projection of neuron activations is inspiring, maybe used in the exploration phase.

- Towards Better Analysis of Deep CNNs (VAST 2016)

  Formulate CNN as DAG and bi-clusters, structural visualization.

Others:

- Model Performance Analysis: Squares (VAST 2016) …
- Interactive Machine Learning…

Design Tasks

High

Priority

1. Presents the learned features of RNNs (hard)
    a) Need to be intuitively **interpretable** (language phenomenon etc.)
    b) Should provide convenient **exploration** interactions (many dimensions of hidden states make no sense)
    c) Experts are interested in **hidden state** and **gate activation** information (haven't validated with experts yet)
2. Comprehensively analyze model's **error/performance**
3. **Comparative** analysis of different models
4. Visualization of RNNs' training process and debugging information

Low

Visual Designs

- Summary View (T2, T4):
  - Error: simple pie chart and data sample for categorical error analysis
  - Training Information: simple line chart for loss curve
- Exploration View (T1.b,c):
  - Exploring interpretable hidden states, salient dimensions of word embedding and **analyze gate activation**.
- Detail View (T2):
  - Use heat map to show detailed hidden state change and gate activations

For T3, add compare (side by side / overlay) to each views

Visual Designs

Exploration View – Semantic Radar

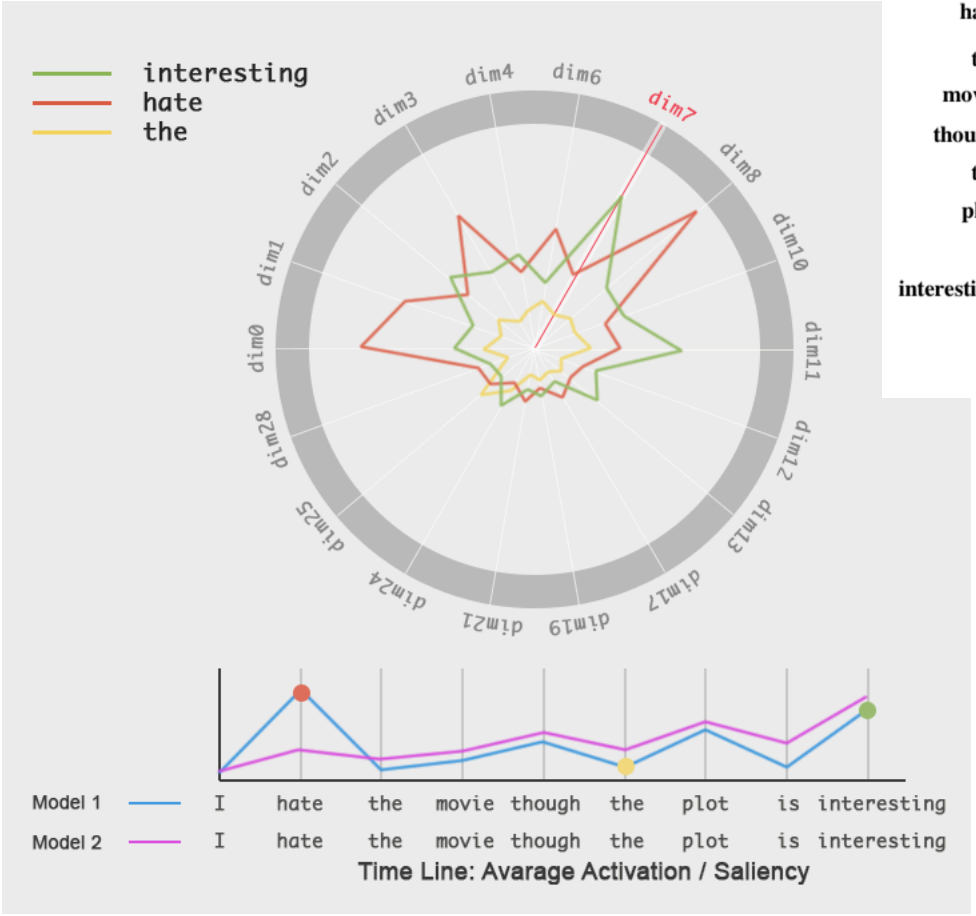

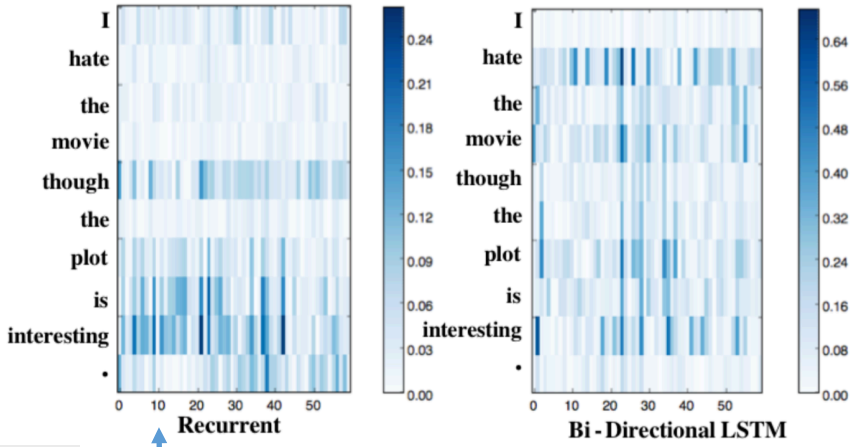

(Li et. al. 2016) Hard to find which dim of word embedding is more salient

Visual Designs

Exploration View – Semantic Radar

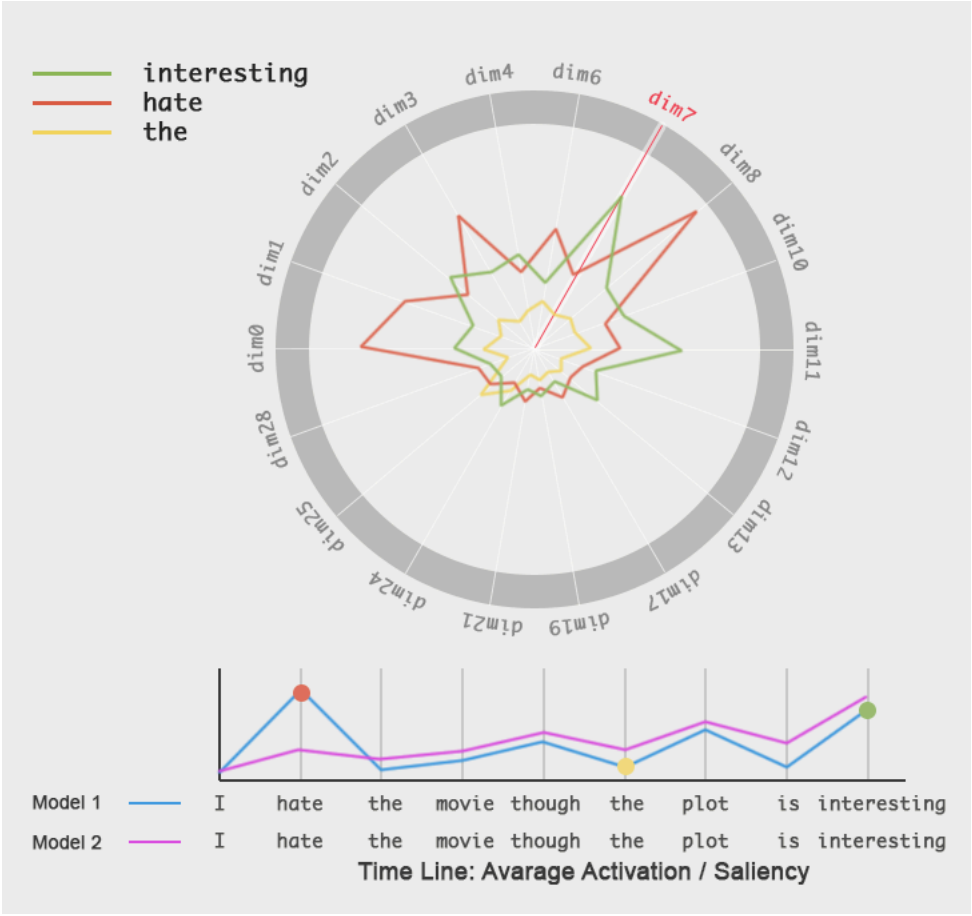

Functions:

- Finding **interpretable** states
- Validate Gate Mechanism
- Compare **keyword** capturing capability

Options:

- **Hidden states** (50~1000d/lr.)
- **Gate Activations** (same as ↑)
- **Saliency scores** of each dim of word embedding (50~1000d)

Visual Designs

Detail View – Heat map (Karpathy et. al. 2015)

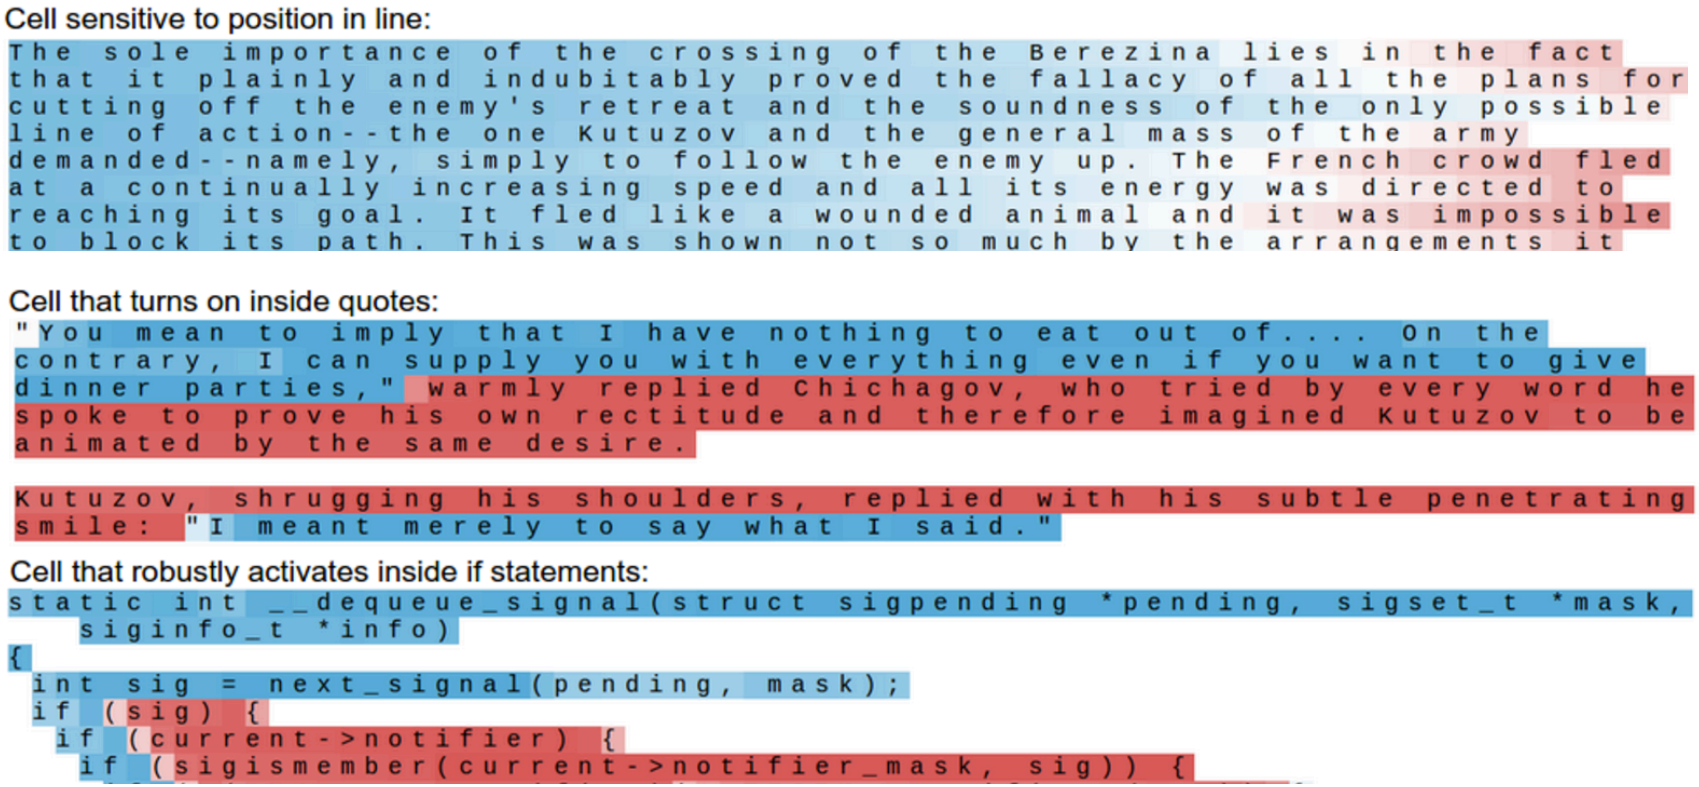

Expected Results

Tools

- RNN implementation: TensorFlow
- Visualization: D3.js
- Framework: Flask (app), Vue.js(front-end)

Datasets:

- Treebanks: Penn Treebank, Stanford Sentiment Treebank
- Plain Text and source code: Linux Kernel Code
- …

Remaining Tasks

Unsolved Major Problems

- Complete and refine visual design
- More literature reviews on existing methods in understanding RNNs and text visualization
- …

Workload:

- Coding
  - RNN and Data Preprocessing ★★★?
  - Visual System ★★★★ ??
- Visual Design and Evaluation ★★★?
- Writing: ★★★

Milestones

	Idea	Coding	Writing
Feb 1st	Fix problem definition & visual designs	Prototype system framework & back-end RNN and data module	Review Related Works
Feb 15th	Refine visual design, compare design alternative	Finish back-end module & Finish Summary View & Prototype Exploration View and Detail View	
Mar 1st		Refine Visualization	Intro, related works & background
Mar 15th		Finish Interactions	System Overview & Visual design
Mar 23rd		Finish experiment and cases	Case study
Mar 29th			Finish all writings